Simpson’s Paradox Explained

statsig.com/perspectives/simpsons-paradox-explained

The Statsig Team

November 2, 2024

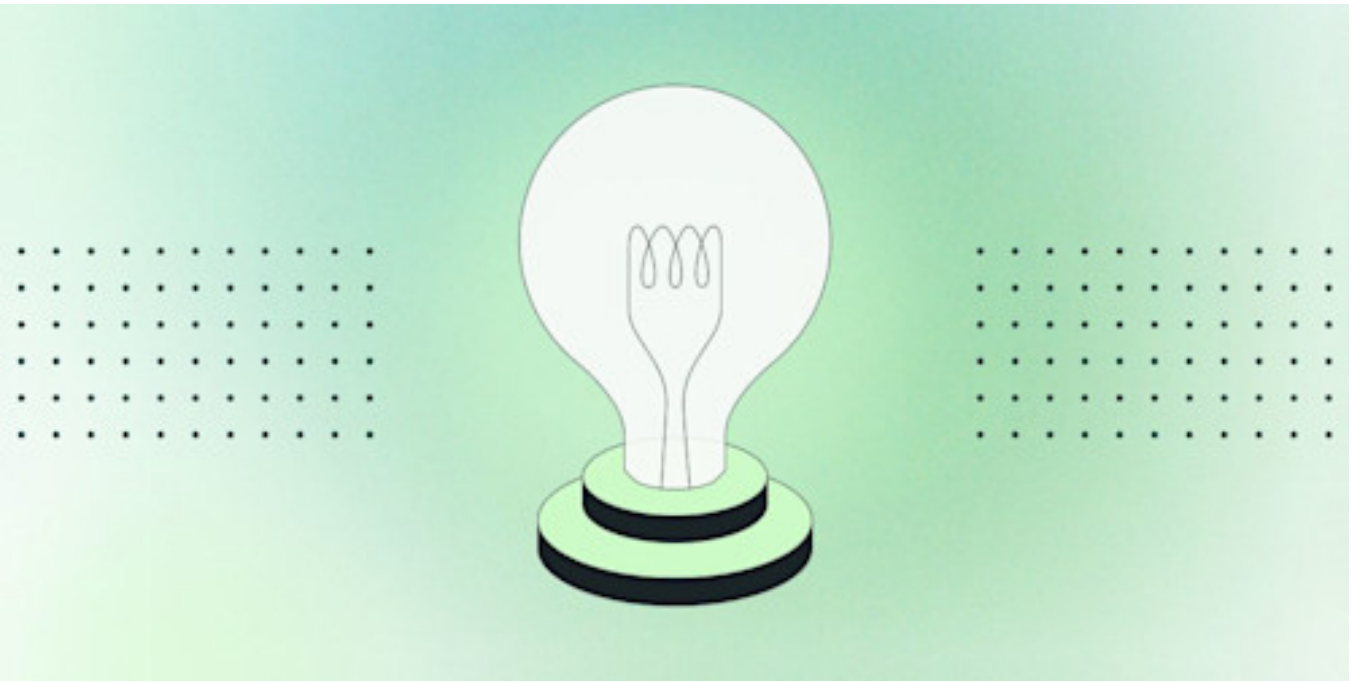

Ever stumbled upon data that tells one story but, when zoomed out, reveals a completely different tale?

Simpson's paradox is that intriguing twist in statistics where trends flip or vanish when you look at the bigger picture. For product managers and engineers, grasping this paradox isn't just academic—it's essential for making sound decisions based on data.

Imagine analyzing user engagement metrics that suggest a feature isn't performing well overall. But wait—dive into specific user segments, and the feature actually shines in each one. How can that be? That's Simpson's paradox in action, and it's a reminder that data can be deceptive if not carefully dissected.

In this blog, we'll explore real-world examples of Simpson's paradox, how it can trip up your A/B testing and experimentation, and strategies to navigate this statistical maze. Let's dive in and unravel this phenomenon together.

Real-world examples illustrating Simpson's paradox

Simpson's paradox emerges in various real-world scenarios, leading to misleading conclusions if not carefully analyzed. In medical studies, treatment effectiveness can appear reversed when data is aggregated across subgroups. For instance, a study might show Treatment A outperforming Treatment B overall, but when separated by gender, Treatment B proves superior for both men and women.

Consider UC Berkeley's admissions data from the 1970s. Initially, it suggested gender bias against women. However, when examining individual departments, no systematic bias was found. The paradox arose because women applied to more competitive departments, skewing the overall results.

During the COVID-19 pandemic, Simpson's paradox was observed in case fatality rates (CFR). Italy had a higher overall CFR than China, but when stratified by age, China's CFR was higher in each age group. The paradox occurred because Italy's population was older, and CFR increases with age.

These examples underscore the importance of considering subgroup data before drawing conclusions from aggregated results. **Simpson's paradox** reminds us that apparent trends can be misleading without proper context and analysis. By recognizing this paradox, researchers and decision-makers can avoid misinterpretations and make more informed judgments based on comprehensive data.

Navigating Simpson's paradox in A/B testing and experimentation

Inconsistent traffic splits in A/B tests can introduce Simpson's paradox, leading to misleading results. For example, if 99% of traffic goes to the control in week one and 50% in week two, the aggregated data may favor the control—even if the variation performs better in both weeks.

To mitigate this, maintain consistent traffic allocation throughout the experiment. If you need to change the split, pause the test, duplicate it, and publish a new one with the desired allocation.

Controlling confounding variables is crucial for accurate experimentation results. Identify potential confounders and use techniques like randomized sampling or blocking to minimize their impact on the primary variables under study.

Be cautious when interpreting **aggregated versus segmented test data**. While overall experiment results should guide decision-making, segment-level metrics can provide insights into potential opportunities. However, avoid basing shipping decisions solely on weak segment-based results; they may be misleading due to Simpson's paradox.

By focusing on comprehensive results, maintaining consistent traffic allocation, controlling confounding variables, and using segmentation to discover possibilities (rather than drive decisions), you can navigate Simpson's paradox effectively. Understanding and addressing this paradox ensures your A/B tests and experiments yield accurate and actionable insights.

Strategies for product owners to mitigate Simpson's paradox in decision-making

Product owners can advocate for **consistent experimental design** to avoid skewed outcomes caused by Simpson's paradox. Changing traffic allocation mid-experiment or reacting to interim results can lead to misleading conclusions. It's crucial to maintain a steady traffic distribution throughout the test.

Encourage analyzing both **segment and overall data** for informed decisions. While segment-level metric interpretations might be misleading, they can guide testing directions and highlight potential opportunities. Focus should remain on aggregate population effects when making shipping decisions.

Promote team education on **statistical nuances** to enhance product success. Understanding concepts like Simpson's paradox, confounding variables, and the sure-thing principle helps teams design better experiments and interpret results accurately. Regular training sessions or workshops can keep everyone up-to-date on best practices.

Consider using tools like Statsig that offer **feature flags and experimentation platforms**. These tools help maintain consistent traffic allocation, manage experiments efficiently, and provide reliable data for decision-making. They can also automate the process of pausing, duplicating, or copying experiments when needed.

By being aware of Simpson's paradox and implementing these strategies, product owners can make data-driven decisions that truly reflect the performance of their features. Careful experiment design, comprehensive data analysis, and a well-informed team are key to navigating the complexities of Simpson's paradox and delivering successful products.

Closing thoughts

Simpson's paradox serves as a powerful reminder that data can be deceptive when not properly analyzed. For product managers and engineers, understanding this paradox is essential for accurate decision-making. By maintaining consistent experimental designs, controlling for confounding variables, and carefully interpreting both aggregated and segmented data, you can avoid the pitfalls of misleading results.

To deepen your understanding of Simpson's paradox and its implications, explore the resources linked throughout this blog. Tools like Statsig can also support your efforts in conducting robust experiments and analyses. **Hopefully, this helps you build your product effectively!**
Interested in learning more? Subscribe to our blogs.

Blogs of interest *

Experimentation Engineering