A recap of the Experimentation Conference 2024 at Booking.com

booking.ai/a-recap-of-the-experimentation-conference-2024-at-booking-com-2f7ad237d95a

Melanie Mueller October 22, 2024

The Machine learning and Data Science blog of Booking.com

By

Melanie Mueller
and
Jorden Lentze

Press enter or click to view image in full size

On May 27, 2024, we organized the first Experimentation Conference at Booking.com in our headquarters in Amsterdam. This invite-only industry event brought together practitioners who run large-scale experimentation programs — as opposed to most other experimentation conferences, which are aimed at people who run experiments. 90 practitioners from 24 companies joined, and discussed common challenges, exchange learnings, and advance experimentation, on 4 specific topics:

- of multiple experiments
- : wrapping code rollouts in experiments
- of running high-quality experiments
- : business outcomes are long term but A/B tests are short term

Each session featured 2 presentations from different companies, followed by a panel discussion on the topic led by industry experts: Ton Wesseling (experimentation evangelist 😀, Kelly Pisane (Booking.com), Lukas Vermeer (Vista Print) and Melanie Kyrklund (Specsavers).

Press enter or click to view image in full size

The Experimentation Conference brought together 90 practitioners running experimentation programs at 24 companies.

Estimation of cumulative effects of multiple experiments

Context: Why are we not seeing the impact of experimentation in our topline revenue? This is a returning question from business leaders. It has been a long term challenge within the industry to understand the real impact of an experiment or a set of experiments on the business.

Join Medium for free to get updates from this writer.

Remember me for faster sign in

ChunWai Kum from Agoda shared how retesting allows to steer experimentation with true impact measurement. Nils Skotara from Booking.com followed up with a roundup of methods we have considered to aggregate the impact from multiple experiments. We saw that different companies approach cumulative impact questions differently and that there's no widely accepted, one-size-fits-all solution.

Press enter or click to view image in full size

Session 1: ChunWai Kum from Agoda (left) and Nils Skotara from Booking.com (right) share different approaches on estimating cumulative effects of multiple experiments

Preventing harm to the business: wrapping code rollouts in experiments

Context: While A/B testing can be used to improve a product, it also acts as a safeguarding mechanism when rolling out new code.

Floor Jubbega and Yftach Zur from Adyen presented how they transformed Adyen's feature flag into an experimentation tool. Daniel Beasley from Vinted explained a new e-values method on how to continuously monitor for harm. This shows that harm prevention can be approached from different angles, depending on the history of experimentation in the company.

Press enter or click to view image in full size

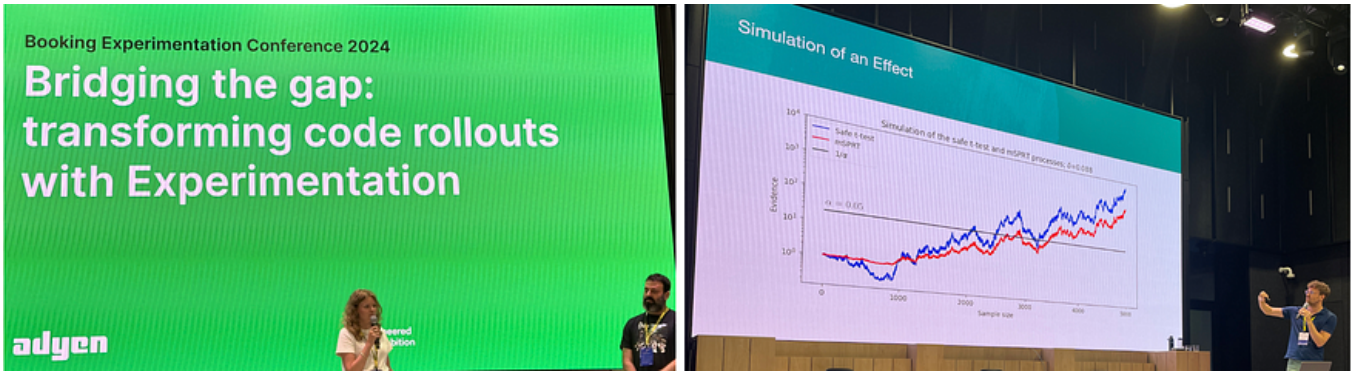

Session 2: Floor Jubbega and Yftach Zur from Adyen (left) and Daniel Beasley from Vinted (right) talk about using experimentation for safe code rollout and monitoring for harm

Creating an experimentation culture

Context: Effective experimentation requires a good experimentation culture in the company: The willingness to run experiments, to run them with high quality, and to listen to the results.

Iryna Horova and Tomasz Kubicki explained how they scaled experimentation culture at the rapidly growing Catawiki, and Ella Johnson-Watts showed what high quality experimentation means at Deliveroo. As a special highlight of the conference, we welcomed back our former Director of Experimentation, Lukas Vermeer. Together with Kevin Anderson from Vista, he discussed building an experimentation culture at Booking.com and at Vista.

This session on experimentation culture proved to be the most popular session, sparking a lot of discussion. Everyone could relate to the challenges of building and maintaining a high-quality impactful experimentation program at their company.

Press enter or click to view image in full size

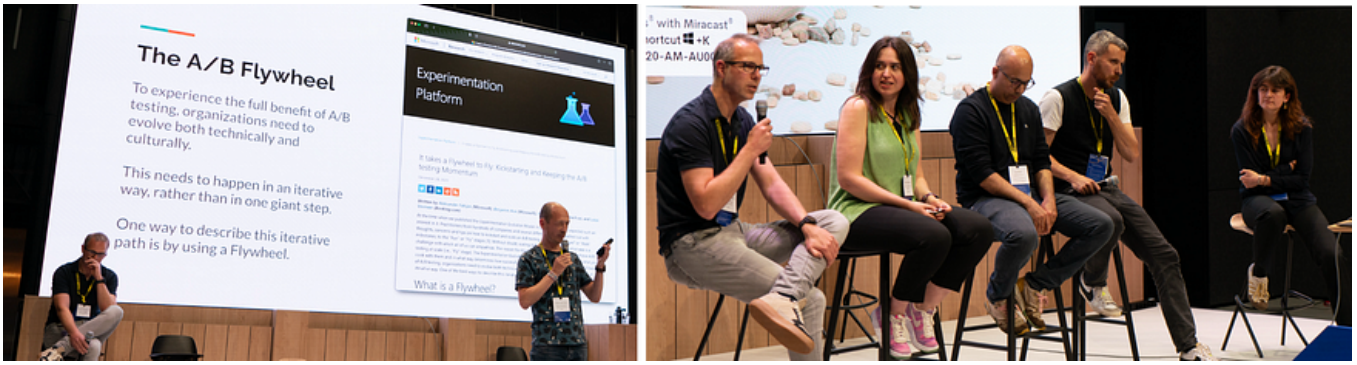

Session 3 on experimentation culture: Lukas Vermeer and Kevin Anderson from Vista (left) gave the conference keynote about building a solid experimentation culture. The topic sparked a lot of discussions in the breaks, and during the panel (right) with Kevin Anderson (Vista), Iryna Horova (Catawiki), Farshad Soleymani (VodafoneZiggo), Alessandro Locatelli (Adevinta), and Ella Johnson-Watts (Deliveroo)

Long-term effects: business outcomes are long term but A/B tests are short term

Context: We are often interested in long-term metrics, such as customer lifetime value, customer retention, etc. But most experiments run only a couple of weeks. How can these "short term" A/B tests be used to inform on long-term metrics?

In the last session on measuring long-term effects, Paresh Nakhe and Geoffrey Decrouez from Zalando explained how to take into account the effect of fashion going out of season in experimentation. Kees Mulder from Bol.com brought out laughs when pondering about getting the wrong answer fast, vs. getting the right answer slow.

Press enter or click to view image in full size

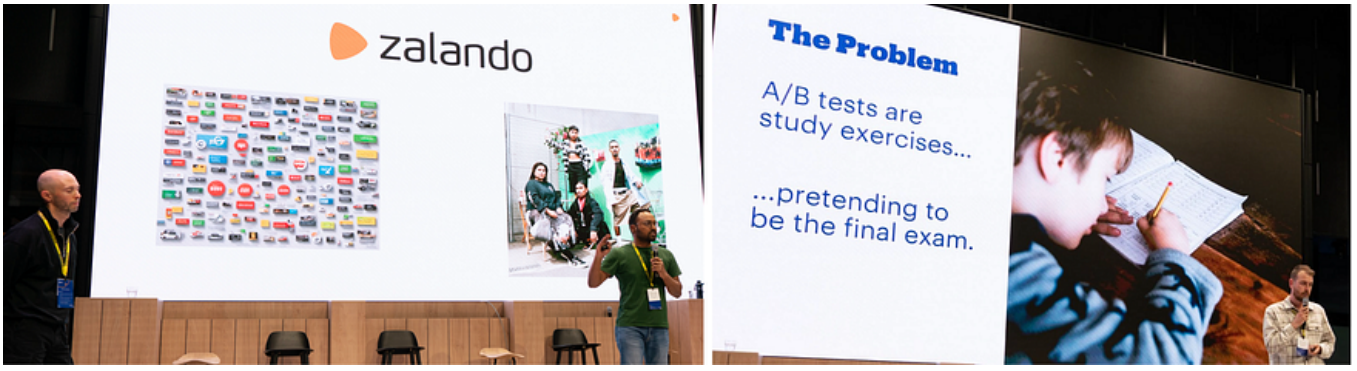

In session 4, Paresh Nakhe and Geoffrey Decrouez from Zalando (left), and Kees Mulder from Bol.com (right) share approaches to estimate long-term effects from short-term A/B tests

We thank all presenters, moderators, and attendees for a great conference! In the after-conference feedback survey, the event received a NPS score of 76%, and 100% of survey respondents said that they would attend again next year 😀. So looking forward to Experimentation Conference 2025 at Booking.com!